

>Beginner Level Task...

Task-1 Iris Flowers Classification ML Project :

This dataset is an Iris dataset usually referred to as the "Hello World" of Machine Learning. The Iris flowers dataset contains numeric attributes, and it is perfect for beginners to learn about supervised ML algorithms, mainly how to load and handle data. Also, since this is a small dataset, it can easily fit in memory without requiring special transformations or scaling capabilities.

Dataset: <http://archive.ics.uci.edu/ml/datasets/iris>
Name: Manish Singh

1.Importing Required Libraries

```
In [26]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as msno
print("All required packages included successfully!")
```

All required packages included successfully!

2.Importing the Dataset

```
In [4]: df = pd.read_csv('D:\Data_Set\Iris.csv')
df
```

| | sepal_length | sepal_width | petal_length | petal_width | species |
|-----|--------------|-------------|--------------|-------------|----------------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 | Iris-setosa |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 | Iris-setosa |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 | Iris-setosa |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 | Iris-setosa |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 | Iris-setosa |
| ... | ... | ... | ... | ... | ... |
| 145 | 6.7 | 3.0 | 5.2 | 2.3 | Iris-virginica |
| 146 | 6.3 | 2.5 | 5.0 | 1.9 | Iris-virginica |
| 147 | 6.5 | 3.0 | 5.2 | 2.0 | Iris-virginica |
| 148 | 6.2 | 3.4 | 5.4 | 2.3 | Iris-virginica |
| 149 | 5.9 | 3.0 | 5.1 | 1.8 | Iris-virginica |

150 rows x 5 columns

3. Data Exploration

```
In [9]: r,c = df.shape
print("Number of rows = ",r)
print("Number of columns = ",c)
```

Number of rows = 150
Number of columns = 5

```
In [5]: df.head()
```

| | sepal_length | sepal_width | petal_length | petal_width | species |
|---|--------------|-------------|--------------|-------------|-------------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 | Iris-setosa |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 | Iris-setosa |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 | Iris-setosa |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 | Iris-setosa |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 | Iris-setosa |

```
In [6]: # to display stats about data
df.describe()
```

| | sepal_length | sepal_width | petal_length | petal_width |
|-------|--------------|-------------|--------------|-------------|
| count | 150.000000 | 150.000000 | 150.000000 | 150.000000 |
| mean | 5.843333 | 3.054000 | 3.758667 | 1.198667 |
| std | 0.828066 | 0.433594 | 1.764420 | 0.763161 |
| min | 4.300000 | 2.800000 | 1.000000 | 0.100000 |
| 25% | 5.100000 | 2.800000 | 1.600000 | 0.300000 |
| 50% | 5.800000 | 3.000000 | 4.350000 | 1.300000 |
| 75% | 6.400000 | 3.300000 | 5.100000 | 1.800000 |
| max | 7.903333 | 4.400000 | 6.900000 | 2.500000 |

```
In [7]: # to display basic data
df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype
--  --
 0   sepal_length  150 non-null     float64
 1   sepal_width   150 non-null     float64
 2   petal_length  150 non-null     float64
 3   petal_width   150 non-null     float64
 4   species       150 non-null     object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

4.Checking Missing Values

```
In [10]: print("Are there any missing values in the dataset ?",df.isnull().values.any())
```

Are there any missing values in the dataset ? False

```
In [14]: msno.bar(df,figsize=(10,6),color='lightgreen')
plt.show()
```



5. Statistical Analysis


```
In [15]: df.describe(include='all').T
```

| | count | unique | top | freq | mean | std | min | 25% | 50% | 75% | max |
|--------------|-------|--------|-------------|------|----------|----------|-----|-----|------|-----|-----|
| sepal_length | 150.0 | NaN | NaN | NaN | 5.843333 | 0.828066 | 4.3 | 5.1 | 5.8 | 6.4 | 7.9 |
| sepal_width | 150.0 | NaN | NaN | NaN | 3.054 | 0.433594 | 2.0 | 2.8 | 3.0 | 3.3 | 4.4 |
| petal_length | 150.0 | NaN | NaN | NaN | 3.758667 | 1.764420 | 1.0 | 1.6 | 4.35 | 5.1 | 6.9 |
| petal_width | 150.0 | NaN | NaN | NaN | 1.198667 | 0.763161 | 0.1 | 0.3 | 1.3 | 1.8 | 2.5 |
| species | 150 | 3 | Iris-setosa | 50 | NaN | NaN | NaN | NaN | NaN | NaN | NaN |

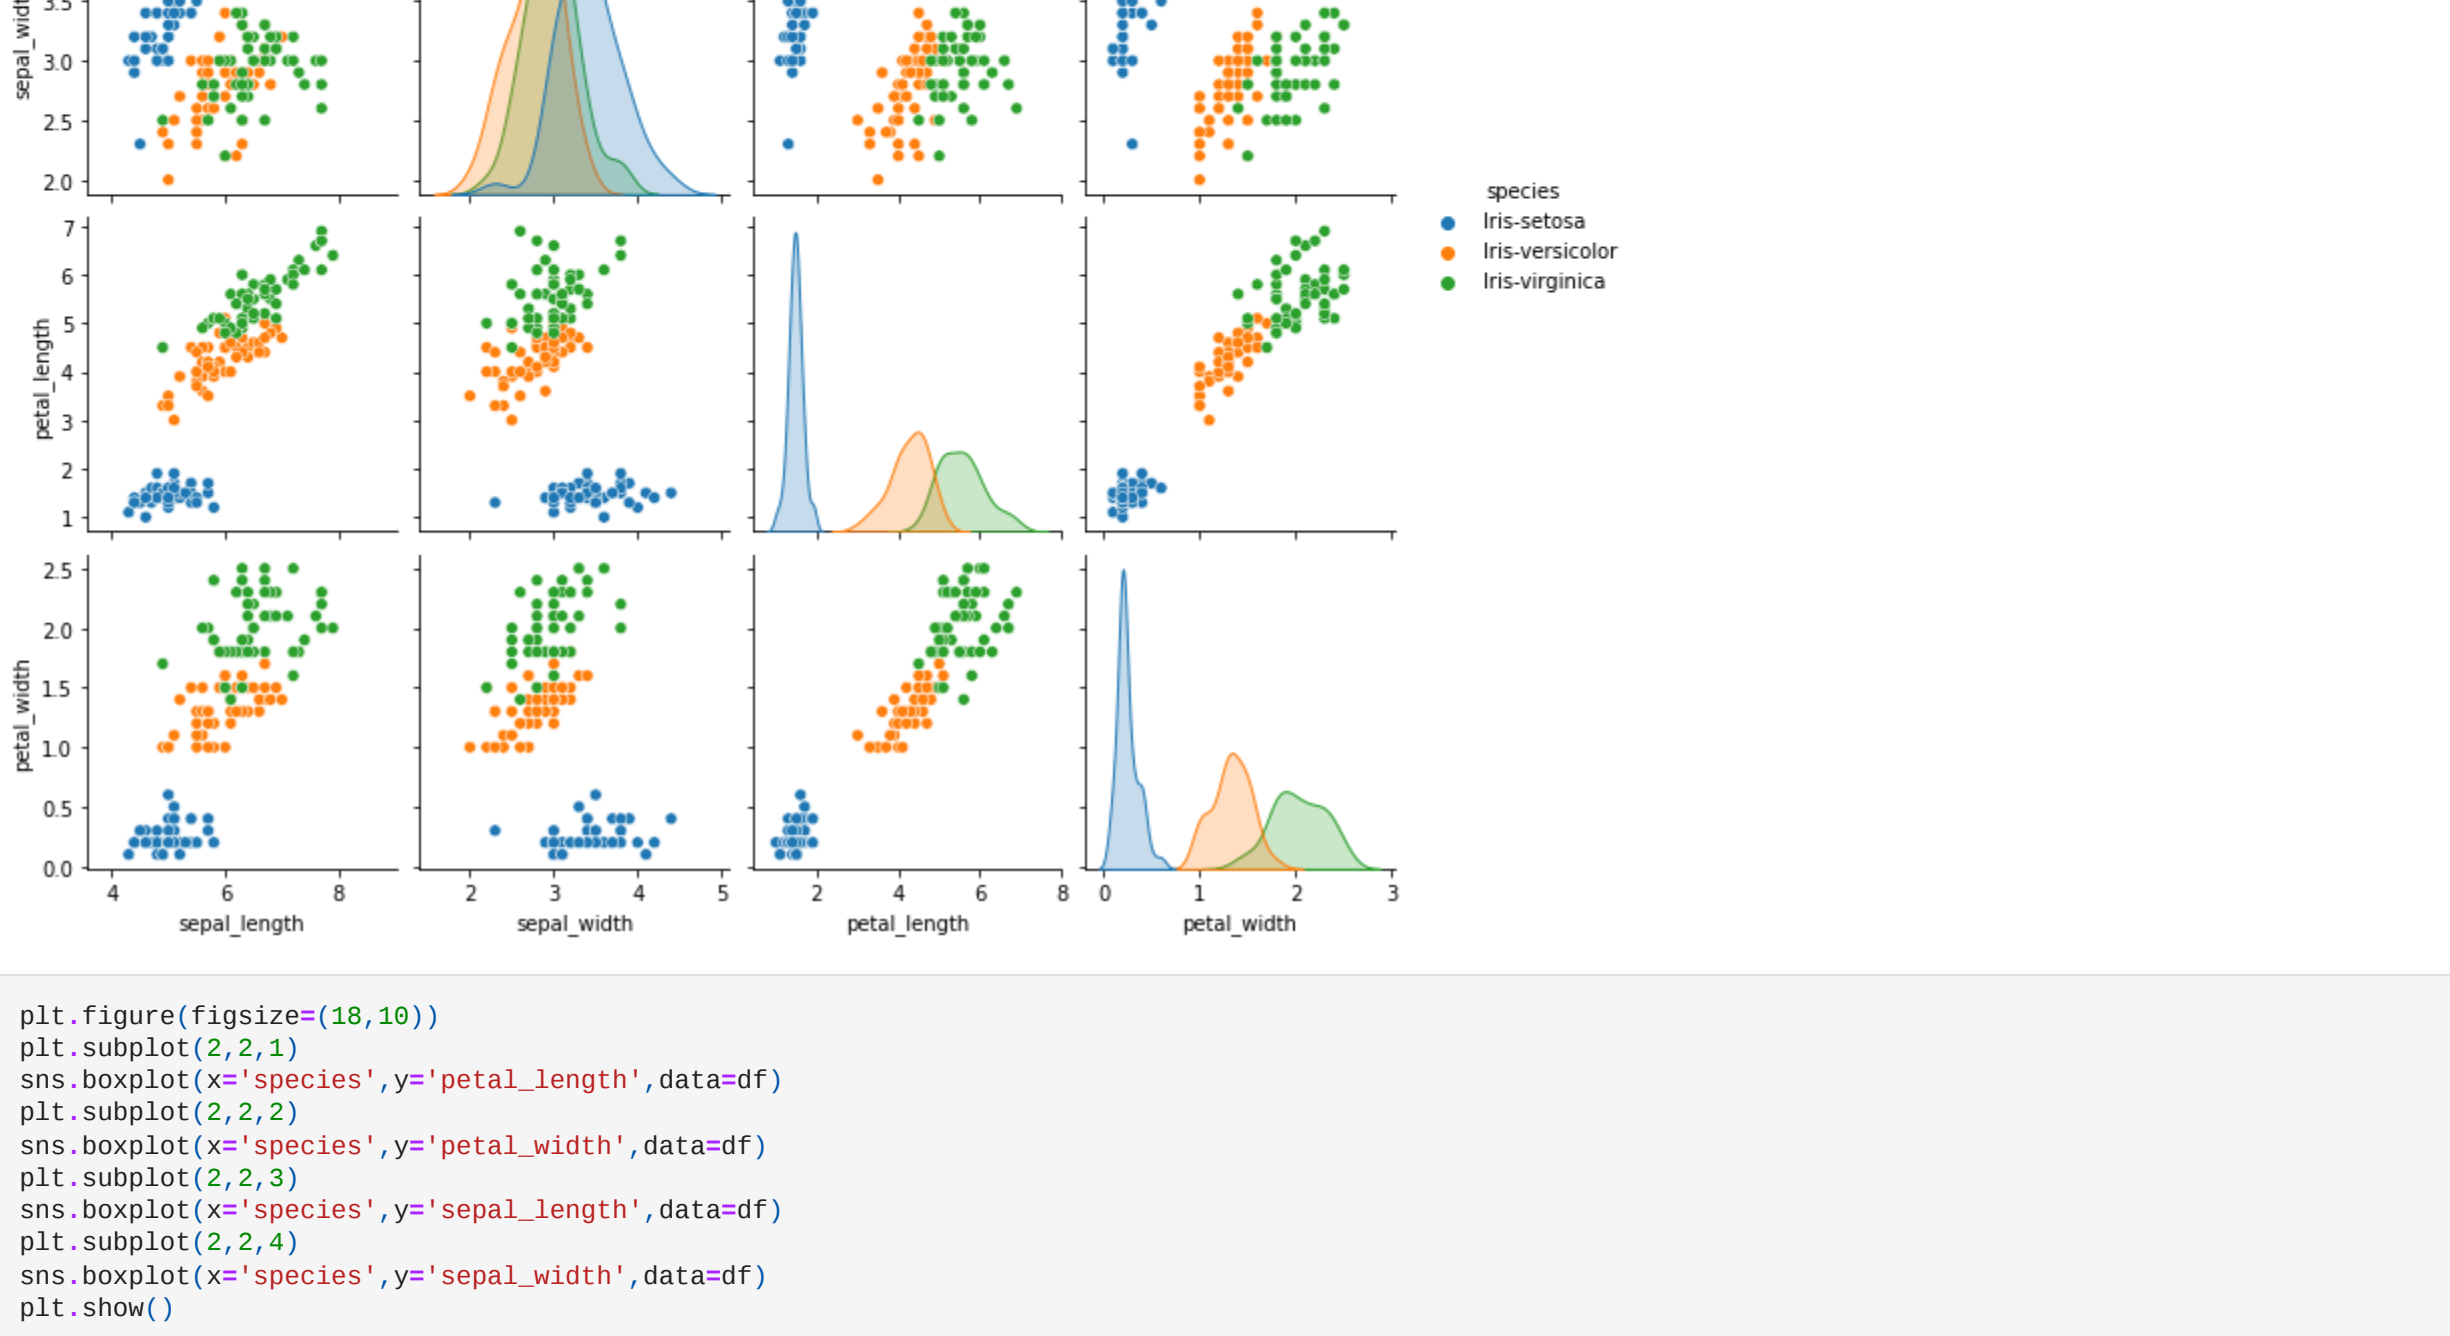
```
In [16]: df['species'].unique()
Out[16]: array(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'], dtype=object)
```

6. Parametric Visualization

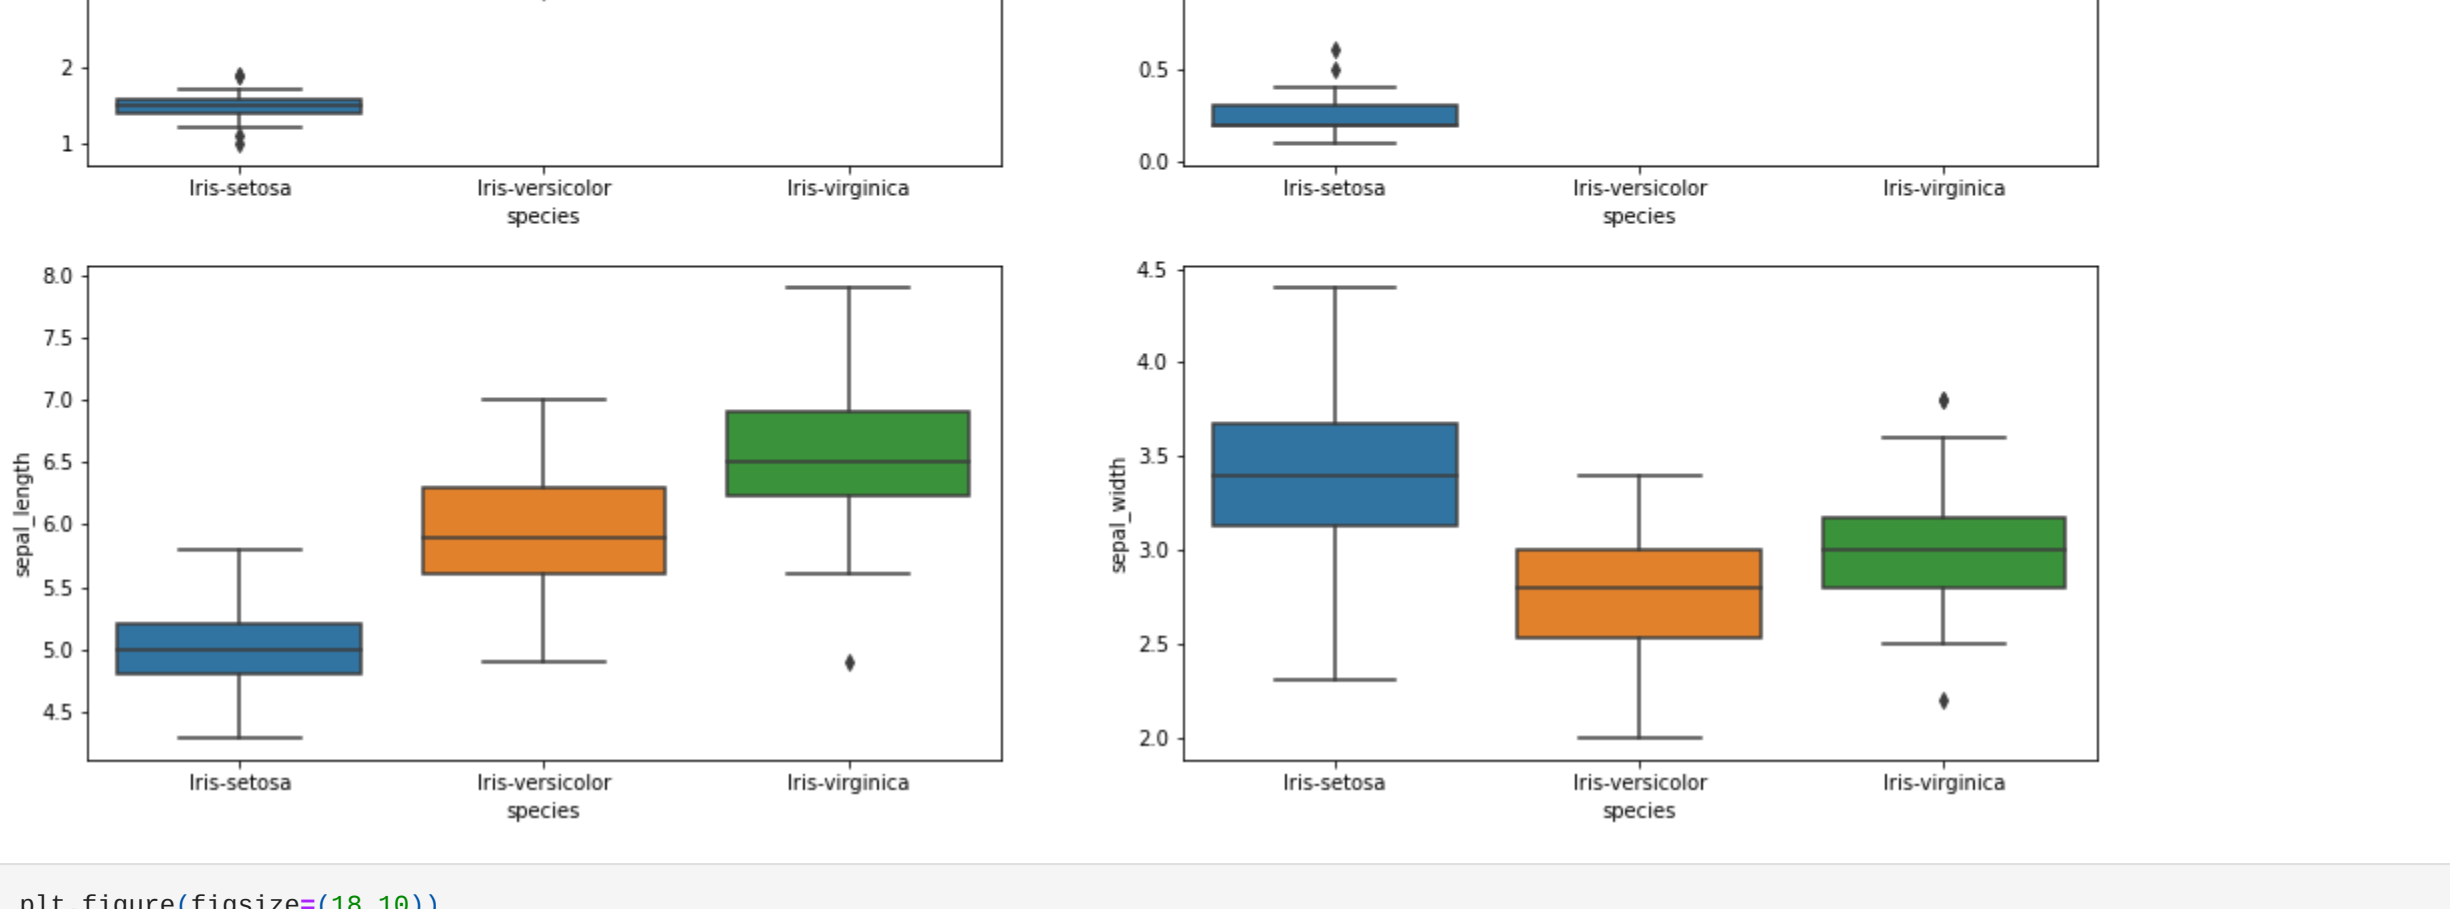
```
In [17]: sns.relplot(x='sepal_length',y='sepal_width',data=df,hue='species',style='species')
g.fig.set_size_inches(10,8)
plt.show()
```



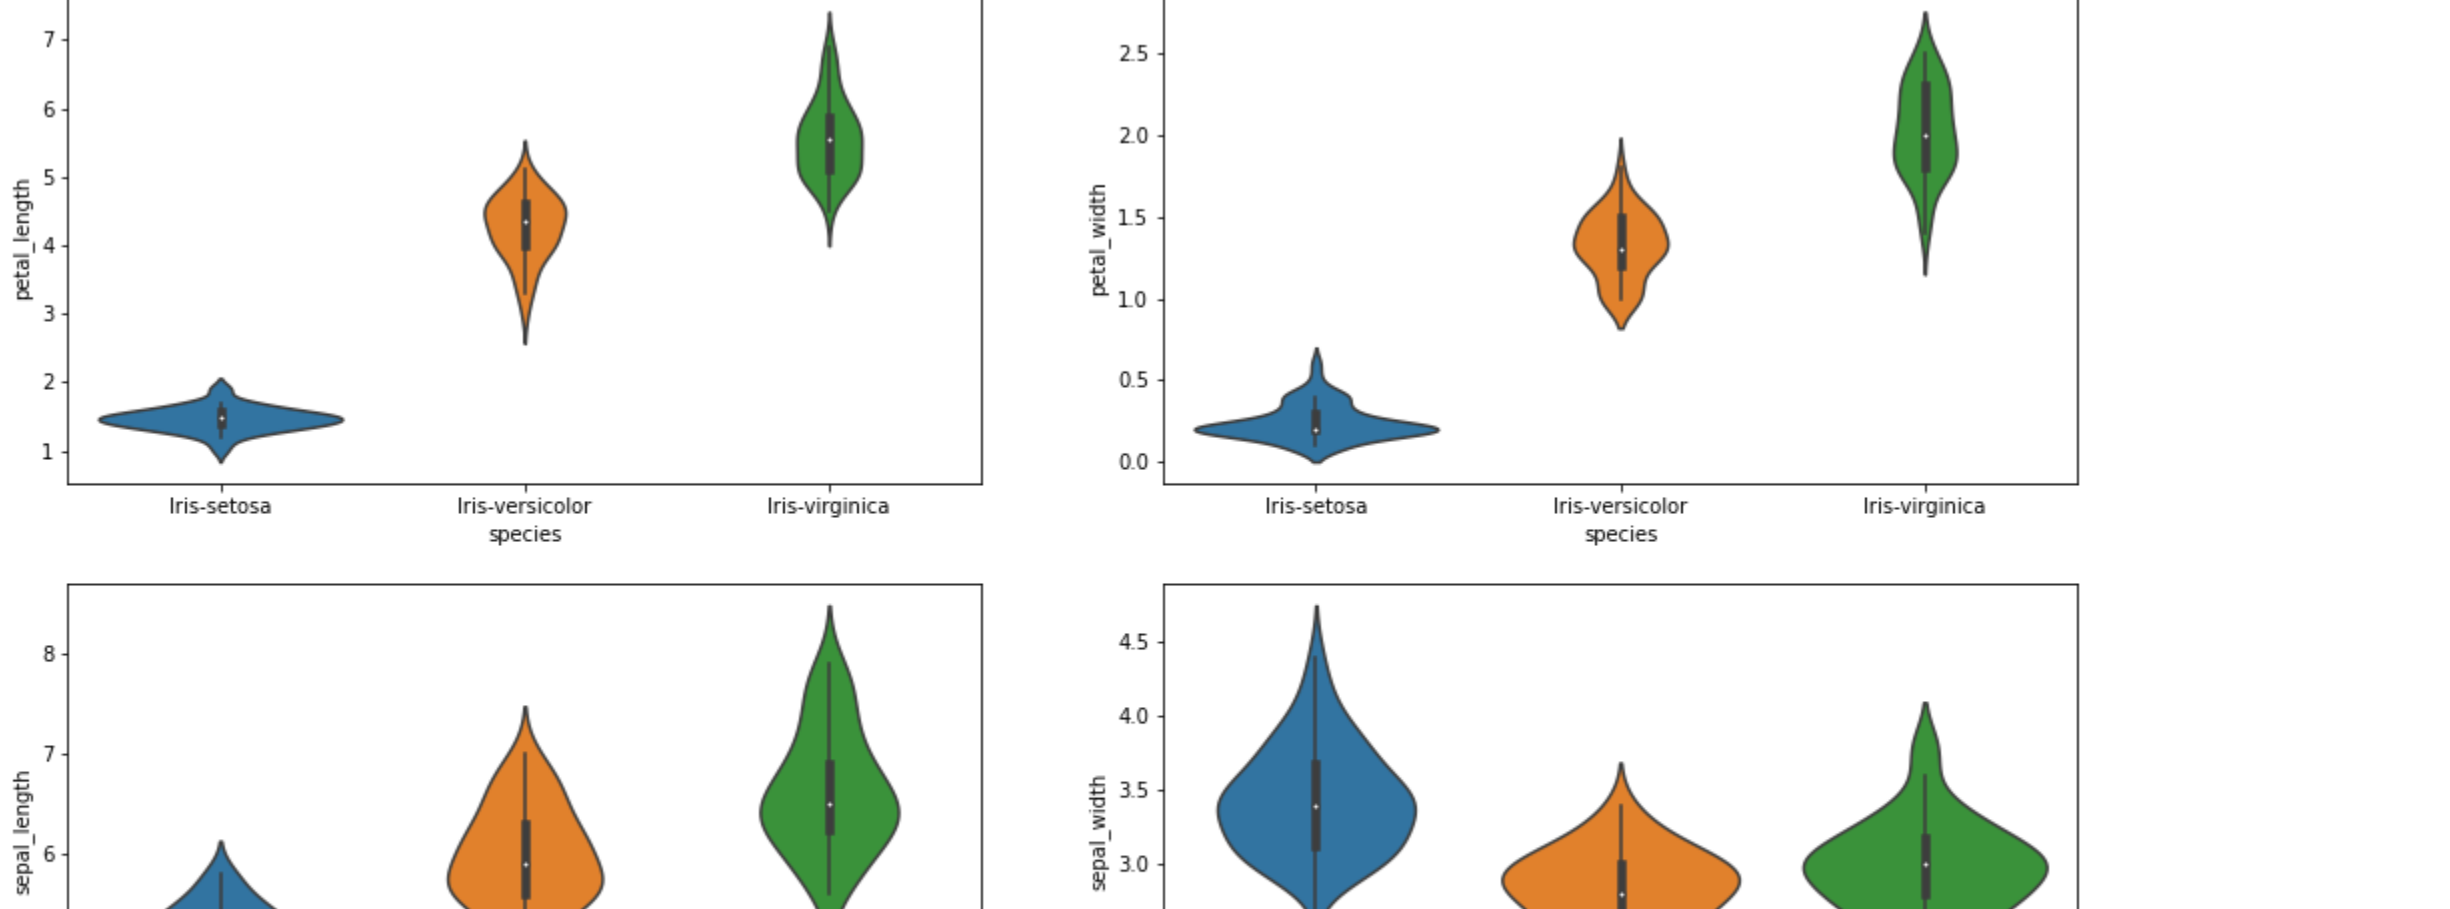
```
In [18]: sns.pairplot(df,hue='species')
plt.show()
```




```
In [19]: plt.figure(figsize=(10,10))
plt.subplot(2,2,1)
sns.boxplot(x='species',y='petal_length',data=df)
plt.subplot(2,2,2)
sns.boxplot(x='species',y='sepal_width',data=df)
plt.subplot(2,2,3)
sns.boxplot(x='species',y='sepal_length',data=df)
plt.subplot(2,2,4)
sns.boxplot(x='species',y='petal_width',data=df)
plt.show()
```




```
In [20]: plt.figure(figsize=(10,10))
plt.subplot(2,2,1)
sns.violinplot(x='species',y='petal_length',data=df)
plt.subplot(2,2,2)
sns.violinplot(x='species',y='sepal_width',data=df)
plt.subplot(2,2,3)
sns.violinplot(x='species',y='sepal_length',data=df)
plt.subplot(2,2,4)
sns.violinplot(x='species',y='petal_width',data=df)
plt.show()
```



```
In [21]: plt.figure(figsize=(10,7))
plt.subplot(2,2,1)
plt.title("Normal distribution of Iris features\n",size=20)
plt.show()
```

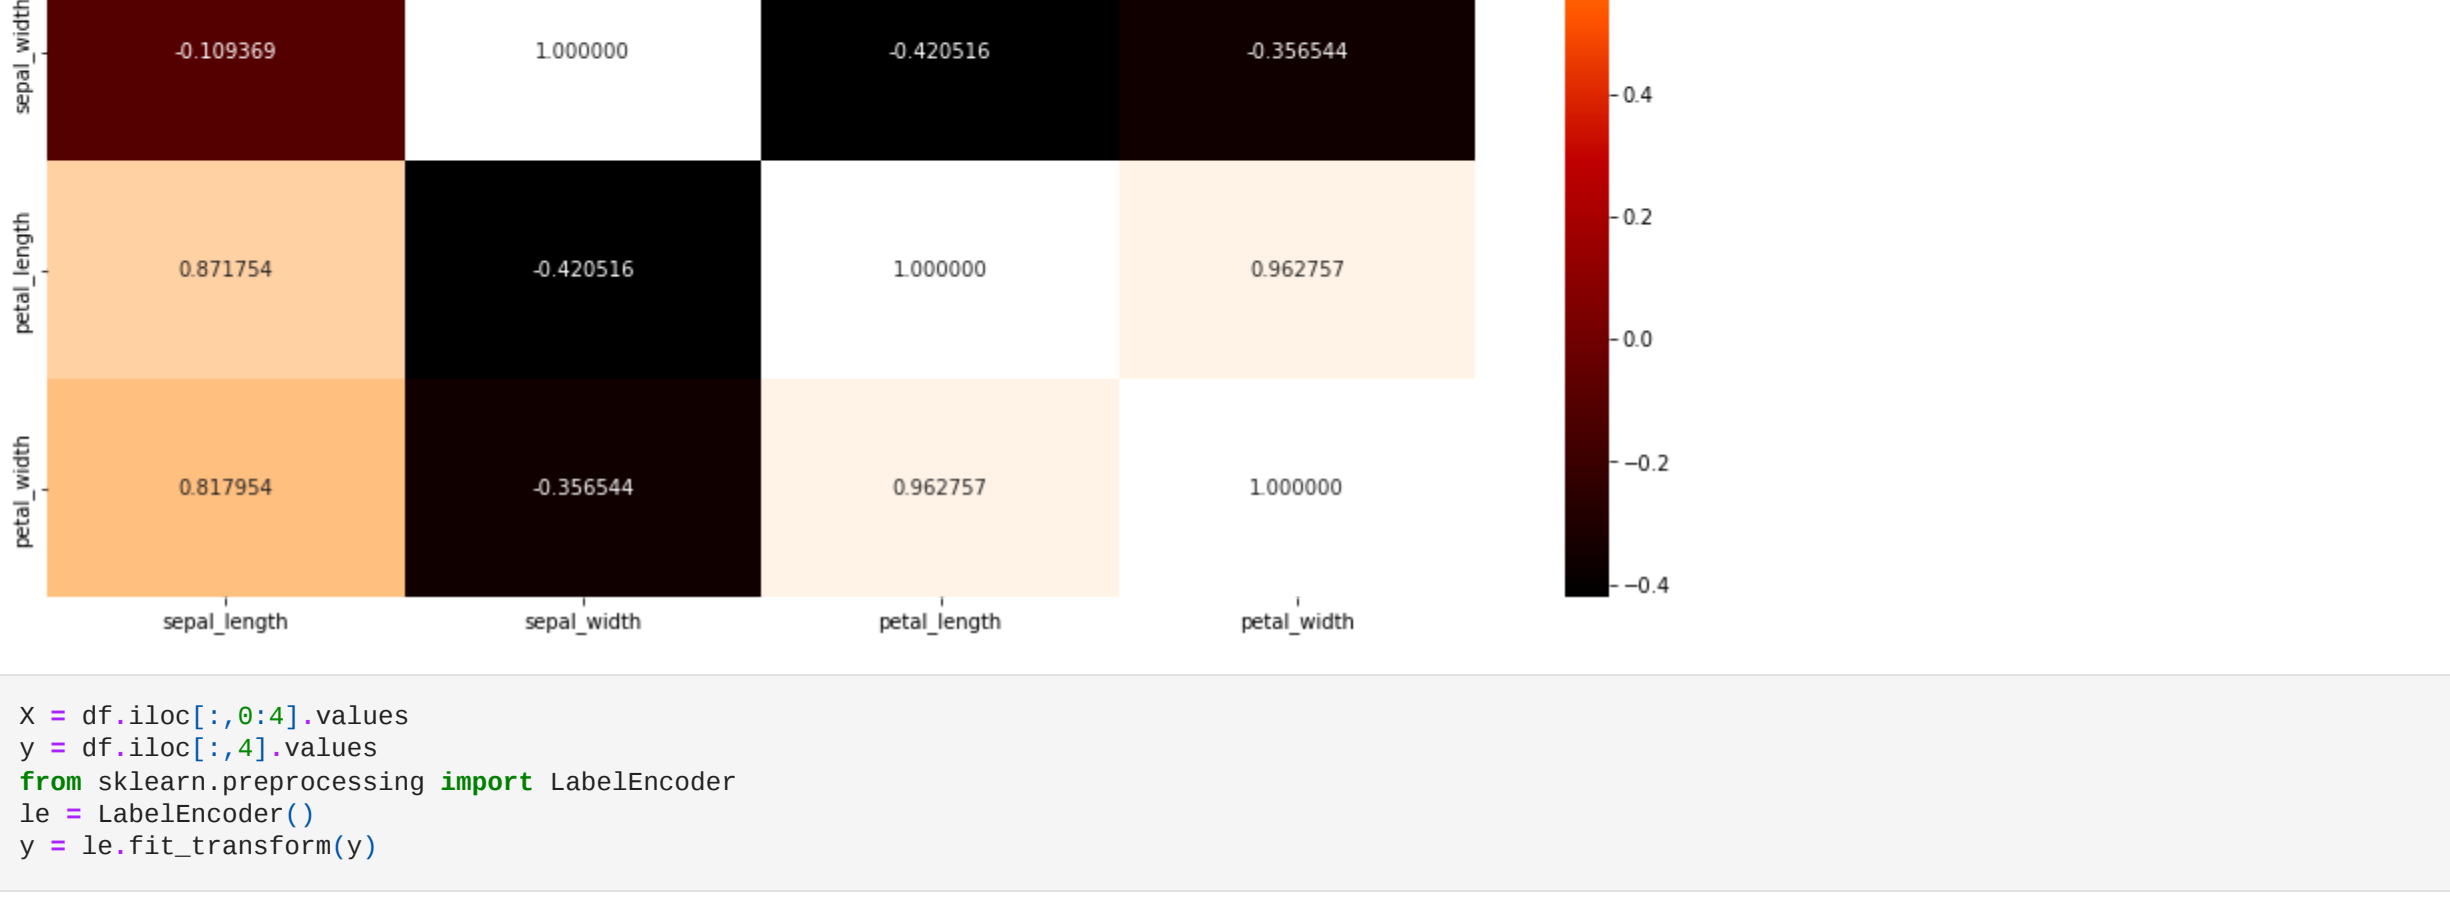


```
In [22]: plt.figure(figsize=(10,7))
sns.violinplot(data=df)
plt.show()
```



7. Attribute Correlation

```
In [23]: plt.figure(figsize=(10,8))
sns.heatmap(corr(df),annot=True,cmap='gist_heat',set_title='Correlation of attributes\n',size=20)
plt.show()
```



```
In [24]: X = df.iloc[:,0:4].values
y = df.iloc[:,4].values
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
y = le.fit_transform(y)
```

8. Metric

```
In [27]: from sklearn.metrics import make_scorer,accuracy_score,precision_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score,precision_score,recall_score,f1_score
print("All necessary metrics included!")
```

All necessary metrics included!

9.Model Selection

```
In [28]: from sklearn.model_selection import KFold,train_test_split,cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import SGDClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC, LinearSVC
from sklearn.naive_bayes import GaussianNB
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=0)
print("All Machine Learning packages included!")
```

All Machine Learning packages included!

10.Random Forest

```
In [29]: rf = RandomForestClassifier(n_estimators=100)
rf.fit(X_train,y_train)
y_pred = rf.predict(X_test)
acc_rf = round(accuracy_score(y_test,y_pred)*100,2)
rf_acc = round(f1_score(X_train,y_train)*100,2)
cm = confusion_matrix(y_test,y_pred)
acc = accuracy_score(y_test,y_pred)
prec = precision_score(y_test,y_pred,average='micro')
recall = recall_score(y_test,y_pred,average='micro')
f1 = f1_score(y_test,y_pred,average='micro')
print("Confusion matrix of Random Forest\n",cm)
print("Accuracy of Random Forest = ",acc)
print("Precision of Random Forest = ",prec)
print("Recall of Random Forest = ",recall)
print("F1 score of Random Forest = ",f1)
```

Confusion matrix of Random Forest

| | [1 0 0] |
|----------|---------|
| [1 0 0] | 15 |
| [0 13 0] | 0 |

Accuracy of Random Forest = 1.0
Precision of Random Forest = 1.0
Recall of Random Forest = 1.0
F1 score of Random Forest = 1.0

11.Logistic Regression

```
In [30]: lg = LogisticRegression(solver='lbfgs',max_iter=400)
lg.fit(X_train,y_train)
y_pred = lg.predict(X_test)
acc_lg = round(accuracy_score(y_test,y_pred)*100,2)
lg_acc = round(f1_score(X_train,y_train)*100,2)
cm = confusion_matrix(y_test,y_pred)
acc = accuracy_score(y_test,y_pred)
prec = precision_score(y_test,y_pred,average='micro')
recall = recall_score(y_test,y_pred,average='micro')
f1 = f1_score(y_test,y_pred,average='micro')
print("Confusion matrix of Logistic Regression\n",cm)
print("Accuracy of Logistic Regression = ",acc)
print("Precision of Logistic Regression = ",prec)
print("Recall of Logistic Regression = ",recall)
print("F1 score of Logistic Regression = ",f1)
```

Confusion matrix of Logistic Regression

| | [11 0 0] |
|----------|----------|
| [11 0 0] | 13 |
| [0 13 0] | 0 |

Accuracy of Logistic Regression = 1.0
Precision of Logistic Regression = 1.0
Recall of Logistic Regression = 1.0
F1 score of Logistic Regression = 1.0

12.K Nearest Neighbours

```
In [31]: knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train,y_train)
y_pred = knn.predict(X_test)
acc_knn = round(accuracy_score(y_test,y_pred)*100,2)
knn_acc = round(f1_score(X_train,y_train)*100,2)
cm = confusion_matrix(y_test,y_pred)
acc = accuracy_score(y_test,y_pred)
prec = precision_score(y_test,y_pred,average='micro')
recall = recall_score(y_test,y_pred,average='micro')
f1 = f1_score(y_test,y_pred,average='micro')
print("Confusion matrix of K Nearest Neighbour\n",cm)
print("Accuracy of K Nearest Neighbour = ",acc)
print("Precision of K Nearest Neighbour = ",prec)
print("Recall of K Nearest Neighbour = ",recall)
print("F1 score of K Nearest Neighbour = ",f1)
```

Confusion matrix of K Nearest Neighbour

| | [11 0 0] |
|----------|----------|
| [11 0 0] | 12 |
| [0 13 0] | 1 |

Accuracy of K Nearest Neighbour = 0.9666666666666667
Precision of K Nearest Neighbour = 0.9666666666666667
Recall of K Nearest Neighbour = 0.9666666666666667
F1 score of K Nearest Neighbour = 0.9666666666666667

13.KNN

```
In [32]: plt.figure(figsize=(20,7))
a_index = list(range(1,50))
a = pd.Series(1)
for i in list(range(1,50)):
    model = KNeighborsClassifier(n_neighbors=1)
    model.fit(X_train,y_train)
    prediction = model.predict(X_test)
    a.append(pd.Series(accuracy_score(y_test,prediction)))
plt.plot(a_index,a,marker='x')
plt.xticks(x)
plt.show()
```

C:\Users\Lenovo\AppData\Local\Temp\ipykernel_6860\41530686925.py:3: DeprecationWarning: The default dtype for empty Series will be 'object' instead of 'float64' in a future version. Specify a dtype explicitly to silence this warning.

a = pd.Series()



14.Gaussian Naive Bayes

```
In [33]: gauss = GaussianNB()
gauss.fit(X_train,y_train)
y_pred = gauss.predict(X_test)
acc_gauss = round(accuracy_score(y_test,y_pred)*100,2)
gauss_acc = round(f1_score(X_train,y_train)*100,2)
cm = confusion_matrix(y_test,y_pred)
acc = accuracy_score(y_test,y_pred)
prec = precision_score(y_test,y_pred,average='micro')
recall = recall_score(y_test,y_pred,average='micro')
f1 = f1_score(y_test,y_pred,average='micro')
print("Confusion matrix of K Nearest Neighbour\n",cm)
print("Accuracy of K Nearest Neighbour = ",acc)
print("Precision of K Nearest Neighbour = ",prec)
print("Recall of K Nearest Neighbour = ",recall)
print("F1 score of K Nearest Neighbour = ",f1)
```

Confusion matrix of K Nearest Neighbour

| | [11 0 0] |
|----------|----------|
| [11 0 0] | 13 |
| [0 13 0] | 0 |

Accuracy of K Nearest Neighbour = 1.0
Precision of K Nearest Neighbour = 0.9666666666666667
Recall of K Nearest Neighbour = 0.9666666666666667
F1 score of K Nearest Neighbour = 0.9666666666666667

14.Linear Support Vector Classifier

```
In [34]: lsvc = LinearSVC(max_iter=4000)
lsvc.fit(X_train,y_train)
y_pred = lsvc.predict(X_test)
acc_lsvc = round(accuracy_score(y_test,y_pred)*100,2)
lsvc_acc = round(f1_score(X_train,y_train)*100,2)
cm = confusion_matrix(y_test,y_pred)
acc = accuracy_score(y_test,y_pred)
prec = precision_score(y_test,y_pred,average='micro')
recall = recall_score(y_test,y_pred,average='micro')
f1 = f1_score(y_test,y_pred,average='micro')
print("Confusion matrix of K Nearest Neighbour\n",cm)
print("Accuracy of K Nearest Neighbour = ",acc)
print("Precision of K Nearest Neighbour = ",prec)
print("Recall of K Nearest Neighbour = ",recall)
print("F1 score of K Nearest Neighbour = ",f1)
```

Confusion matrix of K Nearest Neighbour

| | [11 0 0] |
|----------|----------|
| [11 0 0] | 13 |
| [0 13 0] | 0 |

Accuracy of K Nearest Neighbour = 1.0
Precision of K Nearest Neighbour = 1.0
Recall of K Nearest Neighbour = 1.0
F1 score of K Nearest Neighbour = 1.0

15.Decision Tree Classifier

```
In [35]: dt = DecisionTreeClassifier()
dt.fit(X_train,y_train)
y_pred = dt.predict(X_test)
acc_dt = round(accuracy_score(y_test,y_pred)*100,2)
dt_acc = round(f1_score(X_train,y_train)*100,2)
cm = confusion_matrix(y_test,y_pred)
acc = accuracy_score(y_test,y_pred)
prec = precision_score(y_test,y_pred,average='micro')
recall = recall_score(y_test,y_pred,average='micro')
f1 = f1_score(y_test,y_pred,average='micro')
print("Confusion matrix of K Nearest Neighbour\n",cm)
print("Accuracy of K Nearest Neighbour = ",acc)
print("Precision of K Nearest Neighbour = ",prec)
print("Recall of K Nearest Neighbour = ",recall)
print("F1 score of K Nearest Neighbour = ",f1)
```

Confusion matrix of K Nearest Neighbour

| | [11 0 0] |
|----------|----------|
| [11 0 0] | 13 |
| [0 13 0] | 0 |

Accuracy of K Nearest Neighbour = 1.0
Precision of K Nearest Neighbour = 1.0
Recall of K Nearest Neighbour = 1.0
F1 score of K Nearest Neighbour = 1.0

17.Model Scorer

```
In [36]: res = pd.DataFrame(
    {
        "Model":["knn","Logistic Regression","Random Forest","Naive Bayes","Support Vector Regression","Decision Tree"],
        "score":["acc_knn,acc_lg,acc_rf,acc_gauss,acc_lsvc,acc_dt"],
        "Accuracy_score":["acc_knn,acc_lg,acc_rf,acc_gauss,acc_lsvc,acc_dt,acc"]
    }
)
```

```
res
plt.figure(figsize=(20,8))
ax = sns.barplot(x='Model',y='Accuracy_score',data=res)
labels = (res['Accuracy_score'])
for i,v in enumerate(labels):
    ax.text(1,v+1,align='center',size=15,color='indigo')
```



THANK YOU !